

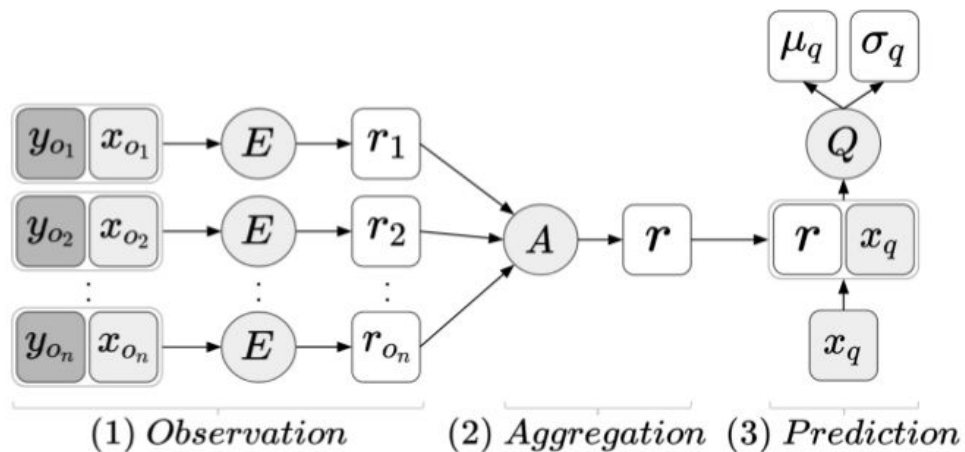
Adaptive Conditional Neural Movement Primitives via Representation Sharing Between Supervised and Reinforcement Learning

M.Tuluhan Akbulut, M.Yunus Seker, Ahmet E. Tekden, Yukie Nagai, Erhan Oztop, Emre Ugur

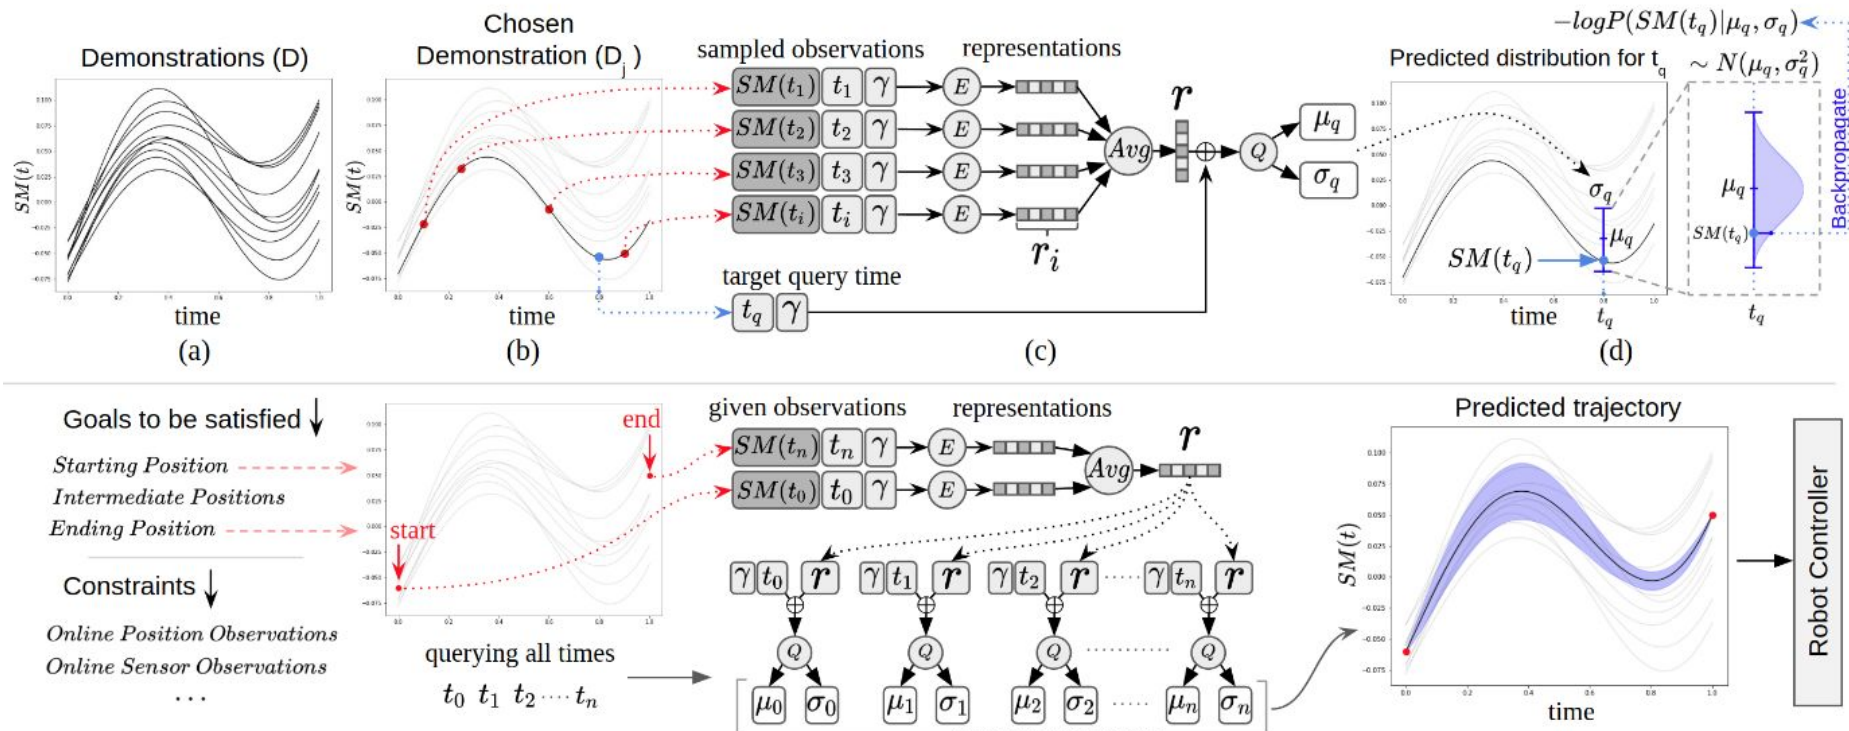


Conditional Neural Processes

- *Conditional Neural Processes*, Garnelo et.al. 2018



Motor primitives based on CNP



Problem and solution

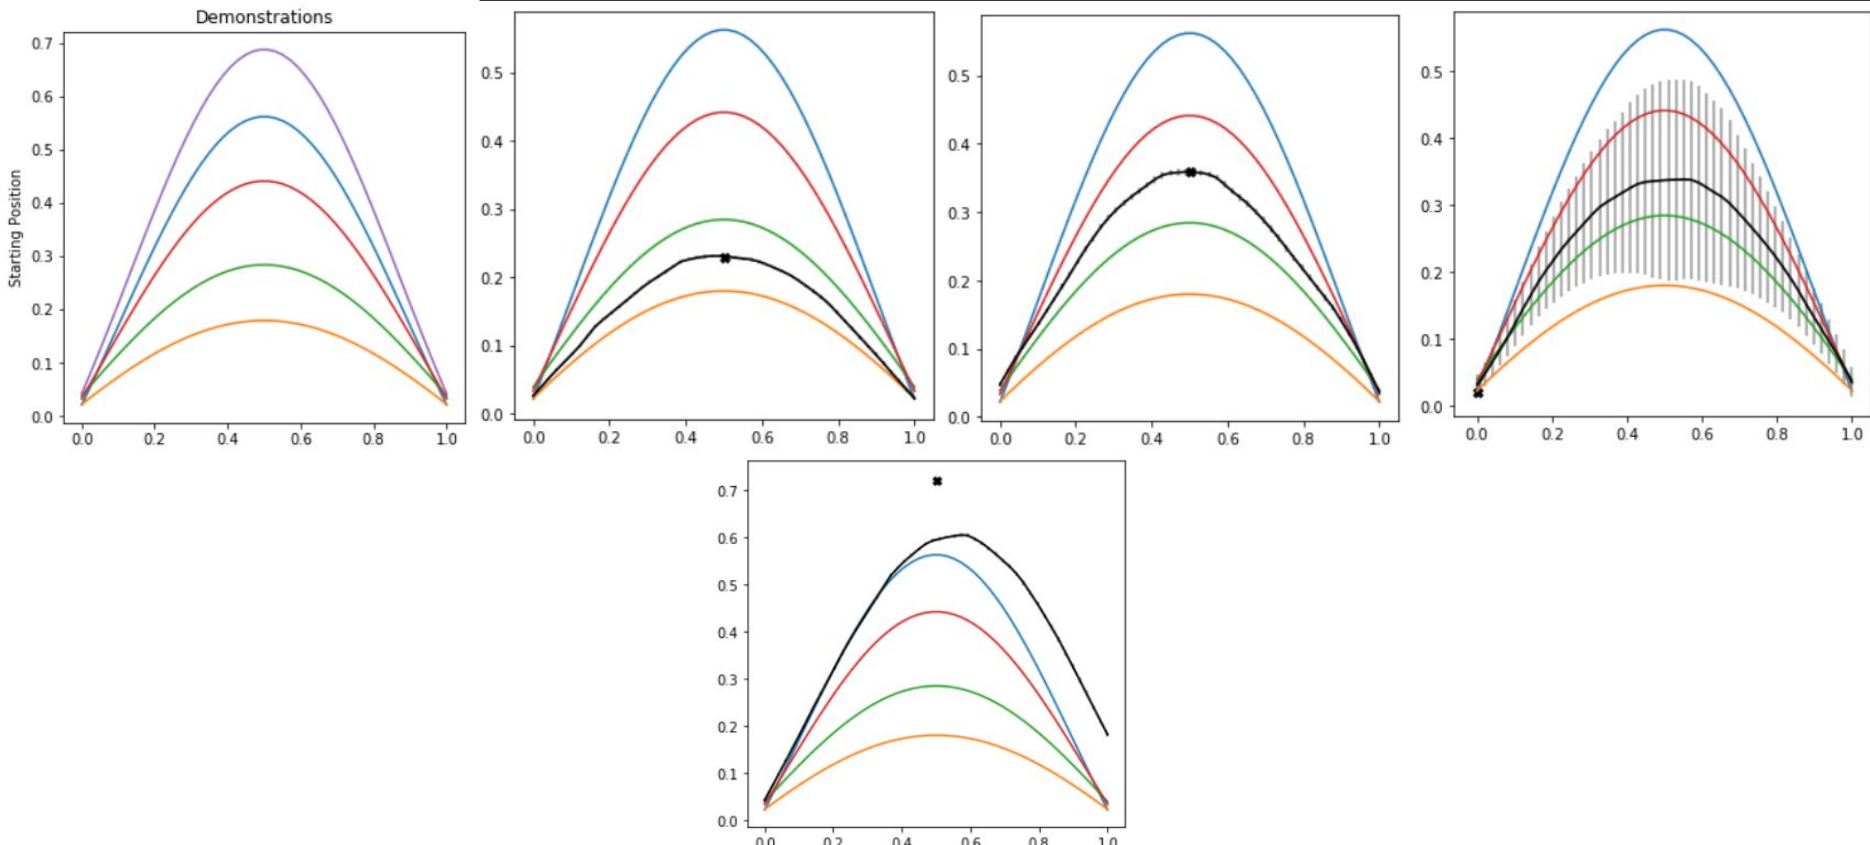
Problems:

- Neural network
- Capabilities are limited with data
- Can't do extrapolation


Solution:

- Self-error detection
- Adaptation algorithm after error
- New task constraints are reached by RL

Capabilities of CNMP



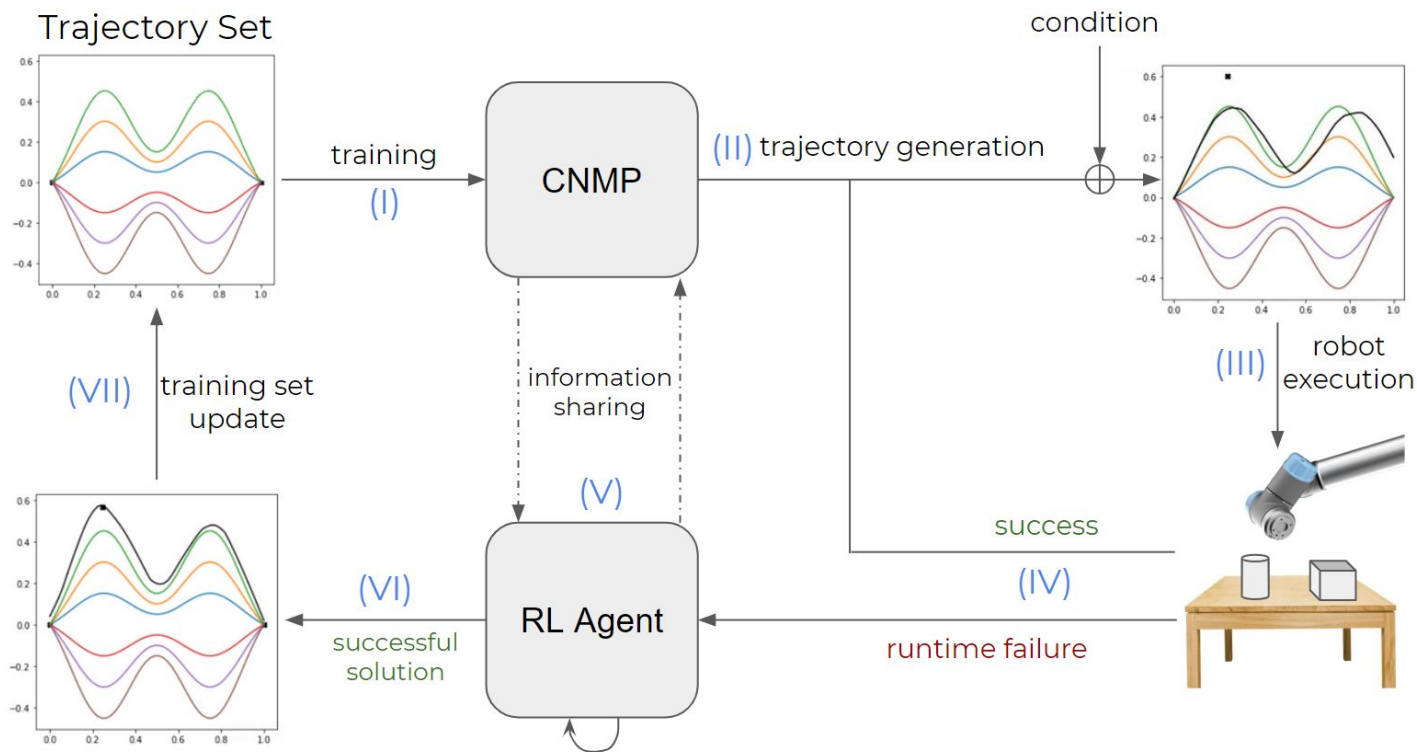
What is RL?



$$\pi_{\theta}(\mathbf{u}_t | \mathbf{x}_t) \rightarrow \tau = \{\mathbf{x}_1, \mathbf{u}_1, \mathbf{x}_2, \mathbf{u}_2, \dots, \mathbf{x}_T, \mathbf{u}_T\}$$


$$\max_{\pi_{\theta}} E_{\tau \sim \pi_{\theta}} \left[\sum_t \gamma^t r(\mathbf{x}_t, \mathbf{u}_t) \right]$$

Method




Correction

- Only reinforcement learning:



- New trajectory should look like expert trajectory (be optimal) to be added CNMP framework.



- RL agent trained on all trajectories (mixed data):



My approach:

policy gradient:
$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) \right) \left(\sum_{t=1}^T r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) \right)$$

maximum likelihood:
$$\nabla_{\theta} J_{\text{ML}}(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) \right)$$



Literature

- How to combine RL and supervised learning:

Most algorithms:



Adding expert demonstrations to replay buffer:

Deep Q-Learning by Demonstrations by *Hester et.al. Google Deepmind, 2018*





$$J_E(Q) = \max_{a \in A} [Q(s, a) + l(a_E, a)] - Q(s, a_E)$$

$$l(a_E, a) = 0 \text{ when } a = a_E$$

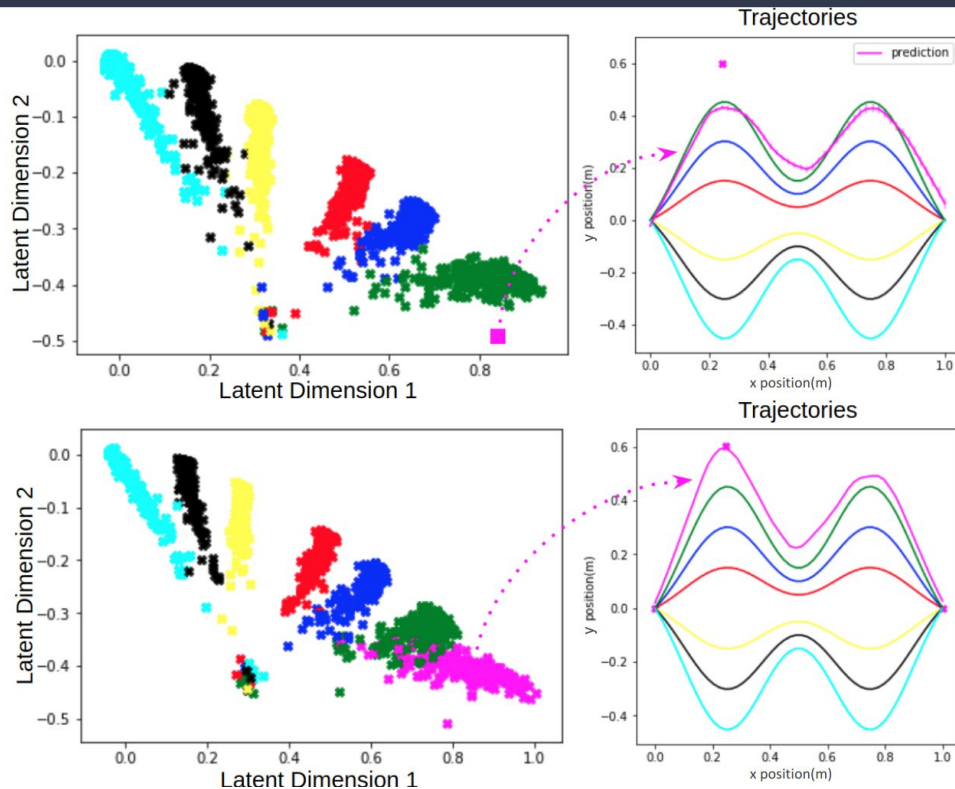
Latent space visualization



Results with a more complex task:

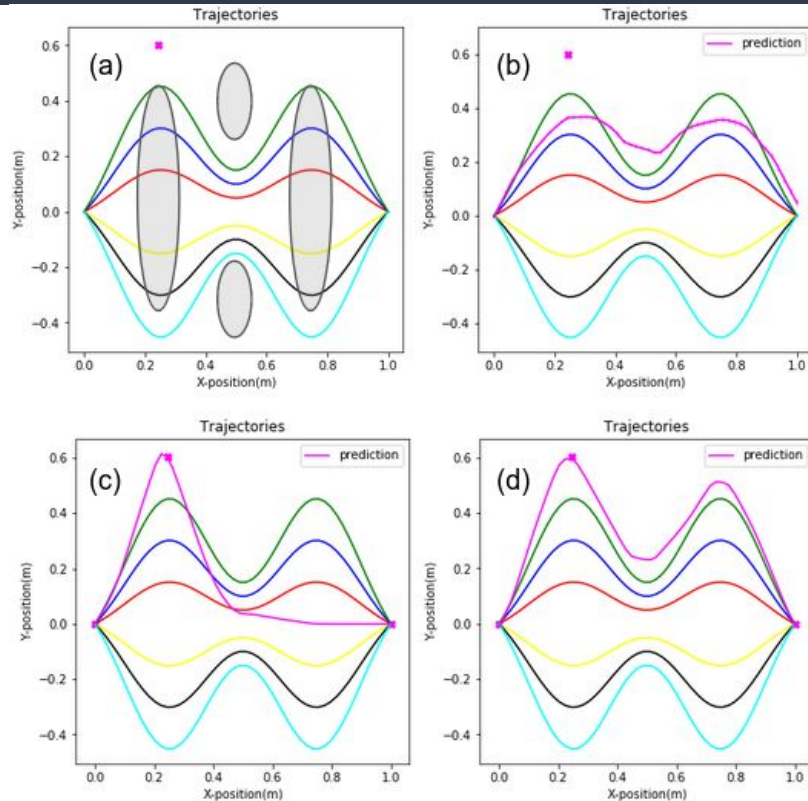


Latent space visualization




Why do we train together

- The policy agent is stochastic
- The gradient is stochastic
- There is no guarantee to preserve the policy shape:







Is it same with regression:



Produced gradients only for 3 points.

In my setup, we let gradient to backpropagate for all time-steps.

Closing the loop:



Experiments:

- Video

Disadvantages

- Since we are training together, it becomes harder to converge(tradeoff between expert trajectory and RL trajectory)
- Harder to differentiate from CNMP
- Takes more iterations, needs optimization.
- Needs further analysis about convergence and representation map
- Oscillation

*Thank You
For Your Atten.*

